

## Pre-Question Data Preprocessing

```
In [1]: ► #import files↔
```

```
In [2]: ► #create a dataframe↔
```

1. **Remove punctuation:** We want to avoid the situation where "maybe" and "maybe?" are considered as different words due to punctuation.
2. **Remove stop words:** Stop words are very likely to appear as top words in each topic, but they don't give us information about what the topic is about. This process will be performed in the CountVectorizer below.
3. **Lemmatization:** Removing inflectional endings and returning the base or dictionary form of a word (i.e. car, cars, car's, cars'  $\Rightarrow$  car) allow having more words related to the topic so he can have a better understanding of it.

## Create Document Term Matrix

When creating the document matrix:

1. I set **max\_df = 0.95**, meaning that I only keep words that show up in at most 95% percent of documents because words that show up in all documents are probably not helpful for identifying topics.
2. I set **min\_df = 2**, meaning that I only keep words that show up in at least 2 documents. I want to have the words that can identify common topics among different documents and ignore those too unique words.
3. I removed stop words, which are not significant and occur frequently.

```
In [4]: from sklearn.feature_extraction.text import CountVectorizer  
cv = CountVectorizer(max_df = 0.95, min_df=2, stop_words='english')
```

```
In [5]: doc_word_matrix = cv.fit_transform(docs['lemmatized'])
```

## LDA

Even though we have 5 topics in the documents, there are lots of subtopics, so I set the number of topics  $k = 10$ . Further exploration on the numbers of topics will be discussed in Question 3. To ensure the result can be reproduced, I set the random state to be 4559. Discussion on random initialization will also be presented in Question 5.

```
In [7]: from sklearn.decomposition import LatentDirichletAllocation  
lda = LatentDirichletAllocation(n_components=10, random_state=4559)
```

Based on the setting, I created two matrixs:

```
In [8]: topic_word_matrix10 = lda.fit(doc_word_matrix)  
doc_topic_matrix10 = lda.fit_transform(doc_word_matrix)
```

## Question 1: Understand Topics

The conventional method that we use to identify a topic is by looking at its top words. However, the number of top words often time has great influence on our interpretation and confidence in the topic.

```
In [12]: > #import packages: math, matplotlib, seaborn↔
```

```
In [13]: > def top_words_for_topic(n, topic_word_matrix, cv, table = True, graph=True)
```

For example, when we look at the top 5 words in each topic, we are likely to recognize that each topic is about:

```
In [14]: > #top 5 words↔
```

```
Out[14]:
```

	topic	top words	possible theme
0	1	[say, year, people, new, report]	report
1	2	[say, use, game, people, user]	game
2	3	[say, mr, labour, party, election]	politics
3	4	[say, year, win, world, set]	game
4	5	[say, company, yukos, oil, mr]	oil business
5	6	[say, game, england, play, player]	sports
6	7	[say, year, market, company, firm]	business
7	8	[say, mr, law, court, police]	law
8	9	[film, good, year, award, star]	film
9	10	[tv, mobile, say, music, digital]	digital entertainment

However, even though when we look at one more word in each topic, we cannot be confident in the statement above. For example, because the word "technology" appears in topic 10, we may consider refine the theme of the topic from "digital entertainment" to "digital technology." Similarly, for topic 1, because of the word "government," we may reconsider this topic as "government-related report."

```
In [67]: > #top 6 words↔
```

```
Out[67]:
```

	topic	top words	possible theme (5)	possible theme (6)
0	1	[say, year, people, new, report, government]	report	government
1	2	[say, use, game, people, user, phone]	game	mobile game
2	3	[say, mr, labour, party, election, blair]	politics	politics
3	4	[say, year, win, world, set, final]	game	game
4	5	[say, company, yukos, oil, mr, plan]	oil business	oil
5	6	[say, game, england, play, player, win]	sports	sports
6	7	[say, year, market, company, firm, price]	business	business
7	8	[say, mr, law, court, police, government]	law	law
8	9	[film, good, year, award, star, win]	film	film
9	10	[tv, mobile, say, music, digital, technology]	digital entertainment	technology

And when we look at the top 10 words in each topic, the topics seems to change again. For example, topic 2 now seems to be about technology rather than mobile games, and topic 10 looks more related to online streaming service.

```
In [16]: ▶ #top 10 words↔
```

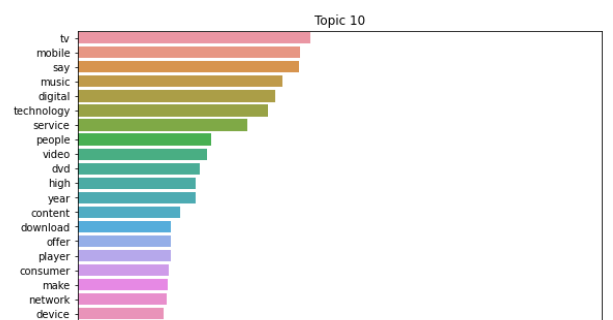









```
Out[16]:
```

topic		top words
0	1	[say, year, people, new, report, government, work, uk, rate, make]
1	2	[say, use, game, people, user, phone, computer, make, software, net]
2	3	[say, mr, labour, party, election, blair, people, government, minister, brown]
3	4	[say, year, win, world, set, final, make, open, time, play]
4	5	[say, company, yukos, oil, mr, plan, russian, firm, court, group]
5	6	[say, game, england, play, player, win, club, time, team, make]
6	7	[say, year, market, company, firm, price, growth, share, rise, bank]
7	8	[say, mr, law, court, police, government, eu, case, right, tell]
8	9	[film, good, year, award, star, win, say, music, include, new]
9	10	[tv, mobile, say, music, digital, technology, service, people, video, dvd]

So the important problem with having to pick a number of words to look at in each topic is that different numbers tell different stories. If we just arbitrarily pick the number without considering the probability distribution of the words, the topic we extract from the words we believe is important can be very misleading. For example, from the barplot below, we can see that word 5 and word 6 and even several words after in each topic have approximately the same probability. Therefore, rather than looking at the top 5 words, we may want to look at several more.

Another potential issue is that we simply take the threshold of number of words and apply it to all topics while different topics have different word distributions. Probably it is more accurate to look at top 10 words for topic 3 because there is a sharp decrease in probability after that while to look at top 20 words for topic 1 because the distribution more even.

```
In [17]: top_words_for_topic(20,topic_word_matrix10,cv,False,True)
```



## Question 2: Examine the Document Topic Probabilities

### Number of Documents in Each Topic

To examine how many documents are about each topic, let's consider the topic of a document to be the topic of highest probability:

```
In [18]: docs['topic'] = doc_topic_matrix10.argmax(axis=1) + 1
```

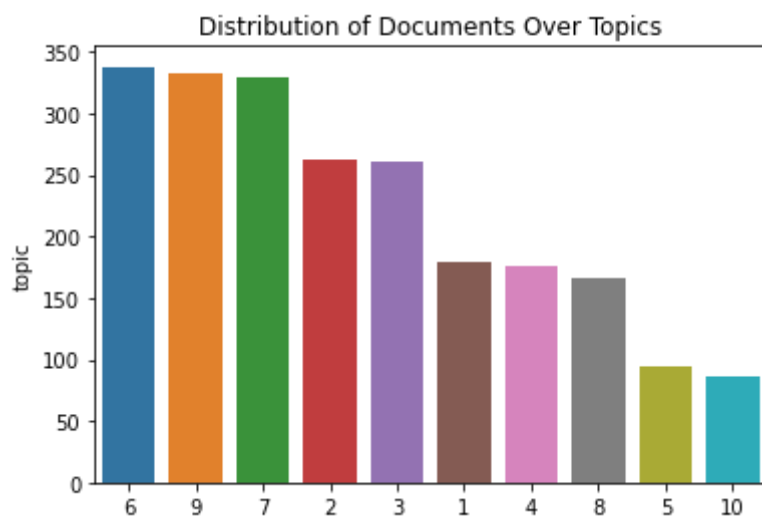
Now look at the number of documents in each topic:

```
In [19]: docs['topic'].value_counts().sort_index()
```

```
Out[19]: 1      179
          2      262
          3      261
          4      176
          5       94
          6      338
          7      330
          8      166
          9      333
         10       86
          Name: topic, dtype: int64
```

```
In [23]: ▶ #Distribution of Documents Over Topics↔
```

```
Out[23]: Text(0.5, 1.0, 'Distribution of Documents Over Topics')
```



We can see that documents are more centered at topic 6(sports), 7(business), and 9(film). One reason can be that these topics are more prevalent among the documents. According to the folder categories, we have more documents in sports and business.

```
In [24]: docs['original folder'].value_counts()
```

```
Out[24]: sport          511
business       510
politics       417
tech           401
entertainment  386
Name: original folder, dtype: int64
```

Now let's see the document distribution within each topic. From the table below, we can observe that topic 6,7,and 9 are centered on 1 specific category.

```
In [25]: docs.groupby(['topic'])['original folder'].value_counts()
```

```
Out[25]: topic  original folder
1      business          71
      politics          59
      tech             35
      entertainment    14
2      tech           255
      business         4
      politics         2
      entertainment    1
3      politics       247
      business        10
      entertainment    2
      sport           2
4      sport         158
      entertainment    13
      tech            4
      politics         1
5      business       61
      entertainment    11
      politics        11
      sport           6
      tech            5
6      sport         335
      tech           2
      politics         1
7      business       319
      tech           5
      entertainment    4
      politics         2
8      politics       92
      business        43
      tech           14
      entertainment    13
      sport           4
9      entertainment   315
      tech            9
      sport           6
      politics         2
      business         1
10     tech           72
      entertainment    13
      business         1
Name: original folder, dtype: int64
```

We can also see that topics with fewer number of documents, such as topic 1 and 8, have a more even distribution of documents. These topics may capture the commonalities in documents from different folders. Are these phenomenons reflected on the top documents in each topic?

## Top Documents in Each Topic



```
In [26]: docs['max topic probabilities'] = doc_topic_matrix10.max(axis=1)
```

```
In [27]: ▶ #find top documents in each topic↔
```

```
Out[27]: topic
1      2022      tech 0.9968633595766595
      2179      tech 0.9968633595766595
      2221      tech 0.9953595903648865
      588      business 0.9947356501994005
      644      business 0.9930759509466004
      1516     politics 0.9402450740733103
      1657     politics 0.9259247708237586
      1498     politics 0.9131727004873222
      1791     politics 0.9130905915365097
      1761     politics 0.8976025796912025
      1861      tech 0.8820769782768075
      417      business 0.8634806172500454
      1422     politics 0.841655275911505
      1446     politics 0.8326128553777089
      1760     politics 0.7978315269989821
      1544     politics 0.7976617077553526
      260     entertainment 0.7971740553247364
      441      business 0.7921457281210003
      808      business 0.7881973936790796
      679      business 0.7826919840420274
8      1710     politics 0.9970680265958527
      1733     politics 0.996938430551284
      1625     politics 0.9964139095028242
      1611     politics 0.9962019500086804
      859      business 0.9957935939202515
      1596     politics 0.9957136143587537
      395      business 0.9954070077721989
      1623     politics 0.995142196621023
      1490     politics 0.9947046661576031
      1420     politics 0.9923709449758934
      1543     politics 0.9898847416027126
      1673     politics 0.9872533632574155
      1521     politics 0.9625724799108706
      1631     politics 0.9560868604191609
      1677     politics 0.9421141463414175
      394      business 0.9387695238434229
      1455     politics 0.9357715060375333
      1489     politics 0.9307987604266335
      1813     politics 0.9268300388908022
      861      business 0.9256876245588085
Name: max topic probabilities, dtype: object
```

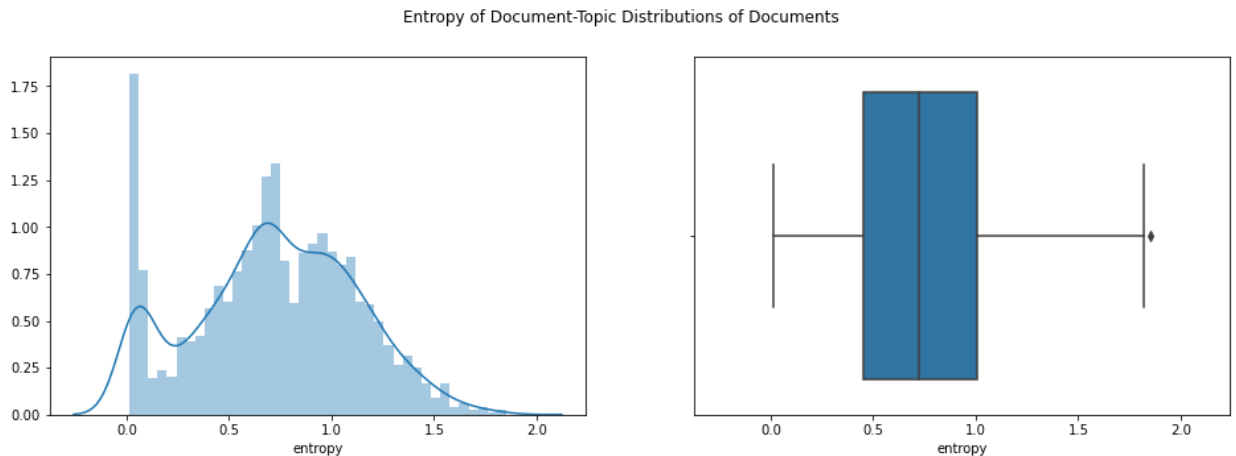
The result is yes. Looking at the top 20 documents in each topic (above), we can discover that topic 1 and topic 8 have documents from different folders while other topics have documents that are consistently from one topic. (For the space reason, only topic 1 and 8 are printed.) Especially, topic 1 have top documents from four different folders.

## Documents that Assign Substantial Weight to More Than One Topic

One good way to measure whether a document spreads weight over multiple topics is **Entropy**. Looking at the histogram and entropy below, we can see that about a quarter of documents have entropy larger than 1, which in this case is considered as high entropy based on the observation in the actual doc-topic probability distributions.

```
In [30]: ▶ #Calculate Entropy↔
```

```
Out[30]: Text(0.5, 0.98, 'Entropy of Document-Topic Distributions of Documents')
```



Comparing the median entropy among different folders, documents from technology and entertainment folder tend to have higher entropy value.

```
In [74]: ▶ #Calculate Entropy for Each Folder↔
```

```
Out[74]: original folder
tech                0.833017
entertainment       0.820671
business            0.791817
politics            0.696227
sport               0.565583
Name: entropy, dtype: float64
```

The high entropy of documents from technology and entertainment is also shown by the fact that these two folders have the highest proportion of documents with entropy greater than 1. It means that technology and entertainment have lots of subtopics that we can identify and explore with details.

```
In [33]: docs.loc[docs['entropy']>=1,'original folder'].value_counts()/docs['origi
```

```
Out[33]: business            0.286275
entertainment            0.357513
politics                 0.242206
sport                   0.113503
tech                    0.324190
Name: original folder, dtype: float64
```

Overall, the higher entropy of a document means that the document spreads weights more evenly

over multiple topics. In the current model, we performed hard assignment and consider the topic of a document to be the topic of highest probability, but the assignment doesn't necessarily grasp the whole picture of the document.

## Question 3: Choice on Number of Topics k

When we categorize documents into different topics, we are exploring the structure of the documents available. With different number of topics, the data structure will also be different. In the question, let's explore 3 different numbers of topics -- 2, 11, and 20 -- and see how the data structure changes compared to the original model with 10 topics by looking at the similarity of top words in pairwise topics.

```
In [34]: ▶ #Build models with k=2,11,20↔
```

```
In [37]: import numpy as np
▶ def count_of_similar_top_words(matrix1, matrix2, k, cv1, cv2):↔
```

### K=2 VS K=10

Based on the count of similar top words, in the model with only 10 topics, each of the topics share relatively many words with one of the topics in the 2-topic model. It can be said that the 10 topics are roughly "divided" into 2 general topics.

```
In [38]: count_of_similar_top_words(topic_word_matrix2, topic_word_matrix10, 20, cv, cv)
```

```
Out[38]:
```

	1	2	3	4	5	6	7	8	9	10
1	8.0	6.0	12.0	3.0	7.0	3.0	7.0	7.0	4.0	4.0
2	6.0	8.0	3.0	9.0	2.0	10.0	2.0	3.0	9.0	8.0

However, looking at the top 20 words of the 2-topic model, we can't really give a label to each topic because they include so many different topics. Therefore, the 2-topic model is much harder to interpret than the 10-topic model.

```
In [39]: top_words_for_topic(20, topic_word_matrix2, cv, True, False)
```

```
Out[39]:
```

	topic	top words
0	1	[say, mr, year, government, new, make, people, labour, company, party, election, plan, tell, minister, uk, country, firm, rise, market, blair]
1	2	[say, year, game, win, make, good, use, people, play, time, new, music, player, world, film, like, come, just, mobile, technology]

### K=11 VS K=10

Comparing the 11-topic model with the 10-topic model, we can see that topic 1 to 10 in both models share many top words along the diagonal of the table, meaning these topics are well preserved, while topic 11 shares relatively few words with the original 10 topics.

```
In [40]: count_of_similar_top_words(topic_word_matrix10,topic_word_matrix11,20,cv,c
```

```
Out[40]:
```

	1	2	3	4	5	6	7	8	9	10	11
1	19.0	4.0	4.0	4.0	2.0	4.0	3.0	5.0	4.0	5.0	5.0
2	4.0	18.0	4.0	2.0	3.0	3.0	2.0	4.0	3.0	6.0	3.0
3	4.0	4.0	20.0	2.0	5.0	2.0	1.0	6.0	2.0	3.0	2.0
4	4.0	2.0	2.0	19.0	1.0	9.0	2.0	2.0	6.0	3.0	5.0
5	3.0	3.0	5.0	2.0	18.0	2.0	5.0	4.0	2.0	2.0	2.0
6	4.0	3.0	2.0	9.0	1.0	19.0	2.0	2.0	6.0	4.0	5.0
7	5.0	2.0	1.0	2.0	3.0	2.0	19.0	1.0	2.0	3.0	2.0
8	5.0	4.0	5.0	2.0	5.0	2.0	1.0	19.0	3.0	3.0	4.0
9	4.0	3.0	2.0	6.0	1.0	6.0	2.0	3.0	14.0	5.0	11.0
10	5.0	5.0	3.0	3.0	1.0	4.0	3.0	2.0	4.0	19.0	4.0

Also, topic 9 in the 10-topic model seems to be decomposed into topic 9 and 11 in the 11-topic model according to corresponding counts. We can also tell this from the top word table below: top words of the old topic 9 related film and music are divided into new topic 9 and 11 by these two categories.

```
In [41]: top_words_for_topic(20,topic_word_matrix10,cv,True,False)
```

```
Out[41]:
```

	topic	top words
0	1	[say, year, people, new, report, government, work, uk, rate, make, rise, figure, child, month, time, job, 000, increase, need, ms]
1	2	[say, use, game, people, user, phone, computer, make, software, net, site, new, technology, mr, firm, online, mobile, pc, microsoft, way]
2	3	[say, mr, labour, party, election, blair, people, government, minister, brown, plan, make, howard, tax, tell, public, prime, leader, tory, chancellor]
3	4	[say, year, win, world, set, final, make, open, time, play, second, good, champion, race, number, title, come, match, beat, break]
4	5	[say, company, yukos, oil, mr, plan, russian, firm, court, group, government, card, claim, buy, state, year, russia, legal, sale, tax]
5	6	[say, game, england, play, player, win, club, time, team, make, ireland, wales, half, year, match, coach, good, come, cup, minute]
6	7	[say, year, market, company, firm, price, growth, share, rise, bank, economy, sale, 2004, china, economic, fall, country, profit, high, month]
7	8	[say, mr, law, court, police, government, eu, case, right, tell, home, lord, rule, make, new, uk, decision, trial, charge, secretary]
8	9	[film, good, year, award, star, win, say, music, include, new, make, actor, song, director, play, band, album, number, chart, british]
9	10	[tv, mobile, say, music, digital, technology, service, people, video, dvd, high, year, content, download, offer, player, consumer, make, network, device]

```
In [42]: top_words_for_topic(20,topic_word_matrix11,cv,True,False)
```

```
Out[42]:
```

	topic	top words
0	1	[say, year, people, new, government, rate, report, work, rise, uk, month, make, figure, child, increase, time, 000, job, high, need]
1	2	[say, use, game, people, phone, user, computer, software, make, net, site, mobile, online, firm, new, mr, technology, pc, mail, internet]
2	3	[say, mr, labour, party, election, blair, minister, people, government, brown, plan, howard, make, tax, tell, prime, public, leader, chancellor, tory]
3	4	[say, win, year, world, set, final, open, make, play, time, second, good, champion, race, beat, title, come, match, break, olympic]
4	5	[say, yukos, company, mr, oil, plan, russian, court, card, firm, claim, government, group, russia, legal, buy, state, case, auction, tax]
5	6	[say, game, england, play, player, win, club, time, team, ireland, make, half, year, match, good, coach, come, rugby, minute, cup]
6	7	[say, year, market, company, firm, share, growth, price, sale, 2004, rise, bank, economy, china, country, economic, profit, fall, analyst, high]
7	8	[say, mr, law, court, government, police, eu, case, tell, right, home, rule, new, lord, make, trial, uk, secretary, charge, minister]
8	9	[film, good, year, award, star, win, say, include, actor, director, new, make, play, prize, actress, role, oscar, movie, tv, british]
9	10	[say, tv, technology, mobile, digital, music, service, people, video, dvd, high, year, content, consumer, network, download, player, make, device, new]
10	11	[music, band, song, album, say, year, good, chart, single, number, make, record, new, wales, artist, radio, rock, uk, release, williams]

Based on this example, we can tell that when we have 10 topics but we choose  $k=11$ , the documents will spread out to occupy 11 topics and have a different configuration. In this case, the original topic 9 becomes new topic 9 and 11 while other topics remain unchanged. However, other documents, due to their nature, may have completely different topics with different models.

## K=20 VS K=10

When we double the number of the topics, it hardly works as we ideally think that each of the original 10 topics is decomposed into 2 subtopics. In this case, we can see from the table below that topic 1-9 in both models share pretty much the same words, especially with original and new topic 3 sharing 19 out of 20 words.

However, several of 20 topics share few words with the original 10 topics, such as the new topic 15 and 16, indicating signs of decomposition and capture of new information.

```
In [43]: count_of_similar_top_words(topic_word_matrix10,topic_word_matrix20,20,cv,c
```

Out[43]:

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
1	9.0	5.0	4.0	5.0	2.0	3.0	2.0	4.0	4.0	3.0	3.0	12.0	2.0	3.0	5.0	1.0	3.0	5.0
2	6.0	14.0	4.0	3.0	3.0	3.0	2.0	4.0	3.0	9.0	2.0	4.0	7.0	2.0	3.0	5.0	7.0	8.0
3	6.0	3.0	19.0	2.0	5.0	2.0	1.0	8.0	2.0	2.0	2.0	5.0	2.0	1.0	2.0	2.0	3.0	3.0
4	3.0	3.0	2.0	10.0	1.0	7.0	3.0	2.0	7.0	1.0	6.0	3.0	2.0	15.0	3.0	1.0	2.0	3.0
5	4.0	1.0	5.0	2.0	15.0	1.0	4.0	5.0	2.0	2.0	2.0	6.0	2.0	2.0	2.0	4.0	3.0	2.0
6	3.0	4.0	2.0	7.0	1.0	17.0	2.0	2.0	6.0	1.0	5.0	3.0	2.0	7.0	3.0	2.0	4.0	3.0
7	2.0	1.0	1.0	2.0	3.0	1.0	14.0	1.0	2.0	2.0	2.0	2.0	2.0	2.0	4.0	3.0	4.0	2.0
8	7.0	2.0	5.0	3.0	4.0	2.0	1.0	15.0	3.0	2.0	2.0	5.0	2.0	1.0	5.0	4.0	4.0	3.0
9	4.0	2.0	2.0	6.0	1.0	5.0	2.0	3.0	14.0	2.0	11.0	4.0	2.0	5.0	4.0	1.0	4.0	3.0
10	4.0	4.0	3.0	3.0	1.0	3.0	2.0	2.0	4.0	4.0	4.0	4.0	2.0	2.0	4.0	1.0	9.0	12.0

Looking at the specific top words in topic 15 and 16, we can see that topic 15 is about broadcasting and topic 16 about crime, topics that are not detected by the 10-word model.

```
In [44]: top_words_for_topic(20,topic_word_matrix20,cv,True,False).iloc[[14,15],]
```

Out[44]:

	topic	top words
14	15	[eu, say, uk, tv, european, country, europe, programme, broadband, straw, year, britain, sky, china, foreign, lift, make, new, constitution, broadcaster]
15	16	[say, mr, attack, fraud, hunt, site, law, patent, net, firm, wales, criminal, company, worldcom, charge, ebbers, mci, security, spam, hunting]

```
In [45]: top_words_for_topic(20,topic_word_matrix10,cv, True,False)
```

```
Out[45]:
```

	topic	top words
0	1	[say, year, people, new, report, government, work, uk, rate, make, rise, figure, child, month, time, job, 000, increase, need, ms]
1	2	[say, use, game, people, user, phone, computer, make, software, net, site, new, technology, mr, firm, online, mobile, pc, microsoft, way]
2	3	[say, mr, labour, party, election, blair, people, government, minister, brown, plan, make, howard, tax, tell, public, prime, leader, tory, chancellor]
3	4	[say, year, win, world, set, final, make, open, time, play, second, good, champion, race, number, title, come, match, beat, break]
4	5	[say, company, yukos, oil, mr, plan, russian, firm, court, group, government, card, claim, buy, state, year, russia, legal, sale, tax]
5	6	[say, game, england, play, player, win, club, time, team, make, ireland, wales, half, year, match, coach, good, come, cup, minute]
6	7	[say, year, market, company, firm, price, growth, share, rise, bank, economy, sale, 2004, china, economic, fall, country, profit, high, month]
7	8	[say, mr, law, court, police, government, eu, case, right, tell, home, lord, rule, make, new, uk, decision, trial, charge, secretary]
8	9	[film, good, year, award, star, win, say, music, include, new, make, actor, song, director, play, band, album, number, chart, british]
9	10	[tv, mobile, say, music, digital, technology, service, people, video, dvd, high, year, content, download, offer, player, consumer, make, network, device]

## Question 4: What happens if you keep the number of topics the same but leave out one of the folders in your analysis?

Topic modeling also highly depends on the data that feeds it. To explore this, I chose to leave out the entertainment folder that have many subtopics and recreated the model.

```
In [50]: > #Build models without the entertainment folder↔
```

Based on the top words of topics in the two models, we can see that the model configuration changed drastically because documents from 4 rather than 5 folders were decomposed to occupy the same number of topics.

While the original topic 4(sports) is preserved as new topic 7 since they share 18/20 words, new topics emerged due to the different decomposition, such as new topic 3(finance) and 4(economy), which share few words with the original 10 topics.



```
In [56]: #row: original topics; column: new topics
count_of_similar_top_words(topic_word_matrix10,topic_word_matrix_no_ent,20)
```

Out[56]:

	1	2	3	4	5	6	7	8	9	10
1	7.0	5.0	4.0	2.0	7.0	8.0	4.0	4.0	4.0	5.0
2	6.0	15.0	5.0	1.0	2.0	6.0	3.0	4.0	3.0	4.0
3	4.0	4.0	3.0	1.0	1.0	5.0	2.0	19.0	2.0	3.0
4	3.0	2.0	3.0	3.0	2.0	3.0	18.0	2.0	7.0	6.0
5	3.0	2.0	6.0	3.0	4.0	8.0	2.0	5.0	2.0	2.0
6	4.0	3.0	3.0	2.0	2.0	3.0	10.0	2.0	12.0	13.0
7	3.0	1.0	6.0	8.0	15.0	5.0	2.0	1.0	2.0	2.0
8	5.0	4.0	7.0	2.0	1.0	7.0	2.0	5.0	2.0	3.0
9	4.0	4.0	4.0	2.0	2.0	4.0	6.0	2.0	6.0	6.0
10	5.0	9.0	3.0	2.0	4.0	4.0	3.0	3.0	4.0	4.0

```
In [55]: top_words_for_topic(20,topic_word_matrix_no_ent,cv_no_ent,True,False)
```

Out[55]:

	topic	top words
0	1	[say, new, sport, make, mac, patient, use, people, uk, domain, team, mr, money, control, add, year, offer, think, johnson, month]
1	2	[say, use, people, mobile, technology, phone, make, service, user, game, computer, music, new, net, mr, digital, network, site, software, work]
2	3	[say, company, firm, year, mr, share, executive, charge, drug, case, new, chief, court, business, deal, financial, profit, bid, make, fraud]
3	4	[say, country, year, world, market, economic, euro, dollar, european, india, company, eu, share, cost, trade, president, deficit, cut, deal, price]
4	5	[say, year, rise, rate, growth, economy, sale, month, 2004, figure, fall, bank, market, price, consumer, firm, quarter, high, increase, economic]
5	6	[say, year, government, mr, company, china, oil, yukos, new, uk, people, make, right, foreign, state, 000, pay, use, country, human]
6	7	[win, game, year, world, say, play, final, time, set, good, title, second, open, champion, match, make, race, beat, come, old]
7	8	[say, mr, government, labour, election, minister, party, blair, people, plan, brown, tell, public, make, howard, issue, prime, leader, tory, tax]
8	9	[say, club, game, play, player, time, chelsea, united, win, good, make, goal, want, league, arsenal, liverpool, cup, team, manager, year]
9	10	[england, say, wales, ireland, win, game, rugby, new, play, try, make, half, france, year, nations, scotland, party, player, coach, time]

## Question 5: Random Initialization

In the original model, the random state is set to 4559. To compare to difference in the topic output due to random initialization, a new model will be produced with a random seed of 2020.

```
In [58]: > #Build models with seed = 2020↔
```

## Among the top words in different topics, are there some topics that remain roughly unchanged?

First of all, we get very a different set of topics from the new model. For example, the first topic in the new model is about technology, while in the original model it is about government.

```
In [59]: top_words_for_topic(10,topic_word_matrix2020,cv,True,False)
```

Out[59]:

topic		top words
0	1	[say, use, people, technology, user, service, net, make, computer, tv]
1	2	[film, good, year, say, award, win, star, include, music, new]
2	3	[say, mr, law, court, case, company, rule, charge, government, new]
3	4	[say, mobile, music, phone, people, make, new, year, use, game]
4	5	[say, mr, labour, party, government, election, people, minister, blair, plan]
5	6	[say, club, game, player, play, win, time, chelsea, make, goal]
6	7	[game, film, play, star, say, year, make, time, new, release]
7	8	[dvd, say, site, high, attack, mini, mac, pc, definition, technology]
8	9	[win, say, year, world, england, play, ireland, final, game, second]
9	10	[say, year, market, rise, company, bank, firm, growth, price, economy]

```
In [60]: top_words_for_topic(10,topic_word_matrix10,cv,True,False)
```

Out[60]:

topic		top words
0	1	[say, year, people, new, report, government, work, uk, rate, make]
1	2	[say, use, game, people, user, phone, computer, make, software, net]
2	3	[say, mr, labour, party, election, blair, people, government, minister, brown]
3	4	[say, year, win, world, set, final, make, open, time, play]
4	5	[say, company, yukos, oil, mr, plan, russian, firm, court, group]
5	6	[say, game, england, play, player, win, club, time, team, make]
6	7	[say, year, market, company, firm, price, growth, share, rise, bank]
7	8	[say, mr, law, court, police, government, eu, case, right, tell]
8	9	[film, good, year, award, star, win, say, music, include, new]
9	10	[tv, mobile, say, music, digital, technology, service, people, video, dvd]

According to the common-word table below, many of the original topics share few common words with any of the new topics, such as topic 1 and 10. However, we can still match some new topics back to the original topics; these topics that remain roughly unchanged :

- (original/column & new/row)
- 3 & 5
- 7 & 10
- 9 & 2

```
In [61]: #row: new topics; column: original topics
count_of_similar_top_words(topic_word_matrix2020,topic_word_matrix10,20,c'
```

Out[61]:

	1	2	3	4	5	6	7	8	9	10
1	5.0	13.0	3.0	2.0	2.0	2.0	2.0	3.0	3.0	7.0
2	4.0	3.0	2.0	6.0	2.0	5.0	2.0	3.0	18.0	5.0
3	4.0	5.0	4.0	2.0	8.0	2.0	4.0	13.0	3.0	2.0
4	6.0	9.0	4.0	3.0	5.0	5.0	5.0	5.0	5.0	8.0
5	5.0	5.0	18.0	2.0	5.0	2.0	1.0	6.0	3.0	3.0
6	3.0	3.0	2.0	6.0	1.0	12.0	1.0	2.0	5.0	3.0
7	6.0	5.0	3.0	7.0	2.0	6.0	2.0	3.0	8.0	5.0
8	3.0	8.0	3.0	3.0	3.0	3.0	3.0	4.0	3.0	7.0
9	4.0	3.0	2.0	14.0	2.0	13.0	2.0	2.0	6.0	4.0
10	6.0	3.0	1.0	2.0	5.0	2.0	18.0	2.0	3.0	3.0

In general, the topic figuration has changed drastically.

## Do the distributions of topics over documents change?

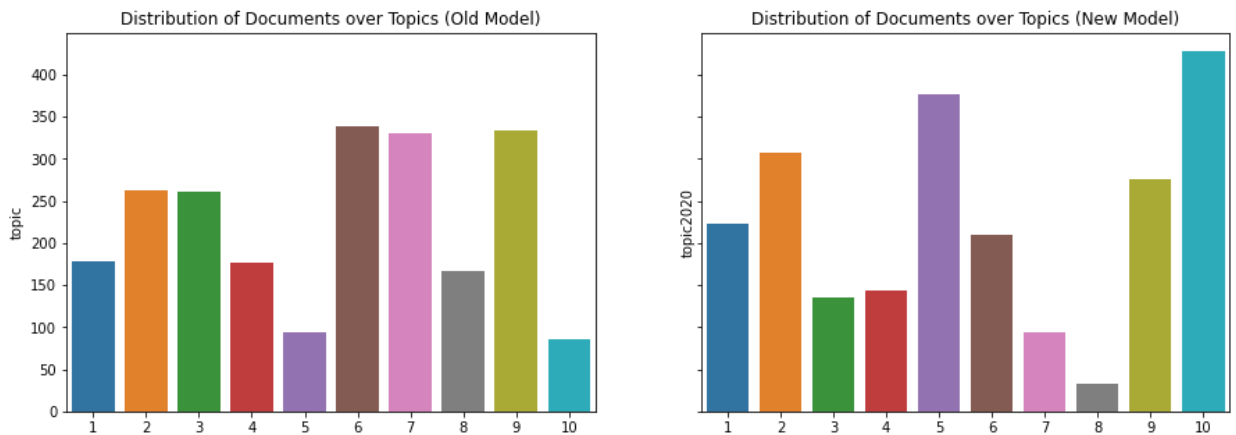
The change in topic configuration is also reflected on the distributions of documents over topics and distributions of topics over documents.

From the barplots below, we can see that the two models have very different distributions of documents over topics. Even the topics like the original 3 and new 5 that match up have different numbers of documents.

```
In [62]: docs['topic2020'] = doc_topic_matrix2020.argmax(axis=1) +1
```

```
In [63]: ▶ #Plot Distribution of Documents over Topics↔
```

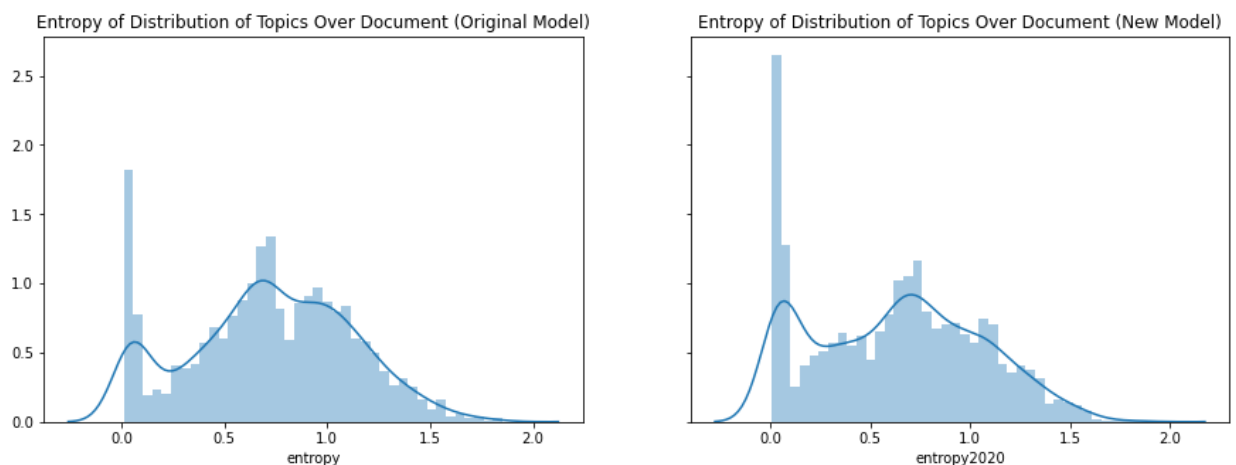
```
Out[63]: Text(0.5, 1.0, 'Distribution of Documents over Topics (Old Model)')
```



To see the change in distribution of topics over documents, we again look at the entropy. Documents in the new model have lower entropy values than those in the original model with a more right-skewed distribution.

```
In [64]: ▶ #Plot Entropy of Distribution of Topics Over Document↔
```

```
Out[64]: Text(0.5, 1.0, 'Entropy of Distribution of Topics Over Document (Original Model)')
```




## Are the documents that were strongly about one topic still about one topic?

Among the documents that originally had a probability of at least 0.8 to assign to 1 topic, only around 61.5% of them are still strong about one topic.

```
In [65]: docs['max topic probabilities2020'] = doc_topic_matrix2020.max(axis=1)
(docs.loc[docs['max topic probabilities'] > 0.8, 'max topic probabilities2020'])

Out[65]: 0.6147859922178989
```

This example shows the instability of model fitting with LDA with random initialization. In reality, the likelihood probability is not smooth but rather bumpy. It is very likely that we stuck on a local maximum and fail to get the best solution. Also, the final space we arrive at also highly depends on where we start the space. That's why with different seeds, the topic models gave different results.



## Question 6 Final Thoughts

Topic models reflect the structure of the data available. There is no ground truth that which model is the right model; rather, topic models are highly unstable. As in our explorations through this homework, topic models depend on many different things: number of top words we choose to interpret the topics, number of topics specified in the model, the data that feeds the model, random initialization, etc.

However, the instability doesn't mean that these models are totally useless. It is just that we need to be very cautious when we use them. We should always use them with human judgements rather than solely rely on them.